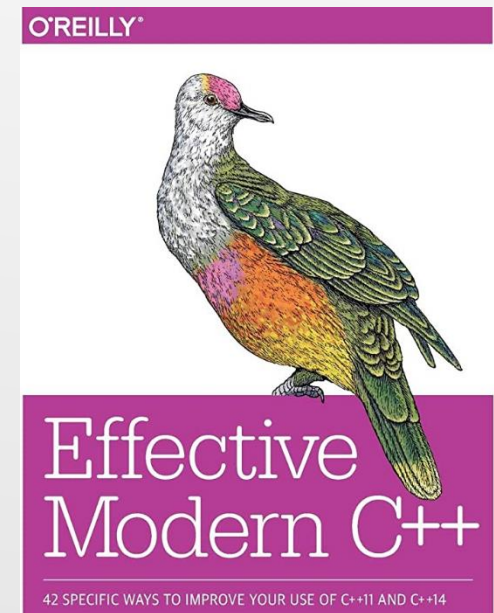


ITEM 16:

Make **const** member functions thread safe

Scott Meyers - Effective Modern C++


Presented By
@7eRoM



- ▶ Consider a polynomial class contains a member function computes the root(s) of a polynomial.

```
class Polynomial {  
public:  
    using RootsType =          // data structure holding values  
        std::vector<double>;   // where polynomial evals to zero  
    ...                         // (see Item 9 for info on "using")  
  
    RootsType roots() const;  
  
    ...  
  
};
```

- ▶ Computing the roots of a polynomial can be expensive, so:
 - ▶ we don't want to do it if we don't have to.
 - ▶ if we do have to do it, we certainly don't want to do it more than once.



```
class Polynomial {
public:
    using RootsType = std::vector<double>;


    RootsType roots() const
    {
        if (!rootsAreValid) {                // if cache not valid

            ...                               // compute roots,
                                           // store them in rootVals

            rootsAreValid = true;
        }

        return rootVals;
    }


private:
    mutable bool rootsAreValid{ false };      // see Item 7 for info
    mutable RootsType rootVals{};            // on initializers
};
```



- ▶ Imagine now that two threads simultaneously call `roots` on a `Polynomial` object:

```
Polynomial p;  
  
...  
  
/*----- Thread 1 ----- */      /*----- Thread 2 ----- */  
  
auto rootsOfP = p.roots();          auto valsGivingZero = p.roots();
```

- ▶ `roots` is a `const` member function, and that means it represents a read operation. Having multiple threads perform a read operation without synchronization is safe.
- ▶ In case of write operation, due to the lack of synchronization, the code has undefined behavior.
- ▶ The problem is that `roots` is declared `const`, but it's not thread safe.



```
class Polynomial {
public:
    using RootsType = std::vector<double>;

    RootsType roots() const
    {
        std::lock_guard<std::mutex> g(m);           // lock mutex

        if (!rootsAreValid) {                     // if cache not valid


            ...                                    // compute/store roots

            rootsAreValid = true;
        }

        return rootVals;

    }                                              // unlock mutex

private:
    mutable std::mutex m;
    mutable bool rootsAreValid{ false };
    mutable RootsType rootVals{};
};
```



- ▶ It's worth noting that because `std::mutex` is a move-only type (i.e., a type that can be moved, but not copied), a side effect of adding `m` to `Polynomial` is that `Polynomial` loses the ability to be copied. It can still be moved, however.
- ▶ In some situations, a `mutex` is overkill.
- ▶ For example, if all you're doing is counting how many times a member function is called, a `std::atomic` counter will often be a less expensive way to go.

```
class Point {                                // 2D point
public:
    ...


    double distanceFromOrigin() const noexcept    // see Item 14
    {                                             // for noexcept

        ++callCount;                            // atomic increment

        return std::sqrt((x * x) + (y * y));
    }

private:
    mutable std::atomic<unsigned> callCount{ 0 };
    double x, y;
};
```

- ▶ Like `std::mutexes`, `std::atomics` are move-only types, so the existence of call `Count` in `Point` means that `Point` is also move-only.
- ▶ Because operations on `std::atomic` variables are often less expensive than `mutex` acquisition and release, you may be tempted to lean on `std::atomics` more heavily than you should.
 - ▶ For example, in a class caching an expensive-to-compute `int`, you might try to use a pair of `std::atomic` variables instead of a `mutex`



```

class Widget {
public:
    ...

    int magicValue() const
    {
        if (cacheValid) return cachedValue;
        else {
            auto val1 = expensiveComputation1();
            auto val2 = expensiveComputation2();
            cachedValue = val1 + val2;           // uh oh, part 1
            cacheValid = true;                   // uh oh, part 2
            return cachedValue;
        }
    }

private:
    mutable std::atomic<bool> cacheValid{ false };
    mutable std::atomic<int> cachedValue;
};

```

- ▶ A thread calls `Widget::magicValue`, sees `cacheValid` as `false`, performs the two expensive computations, and assigns their sum to `cachedValue`.
- ▶ At that point, a second thread calls `Widget::magicValue`, also sees `cacheValid` as `false`, and thus carries out the same expensive computations that the first thread has just finished. (This “second thread” may in fact be several other threads.)


```
class Widget {
public:
    ...

    int magicValue() const
    {
        if (cacheValid) return cachedValue;
        else {
            auto val1 = expensiveComputation1();
            auto val2 = expensiveComputation2();
            cacheValid = true;                // uh oh, part 1
            return cachedValue = val1 + val2; // uh oh, part 2
        }
    }

    ...

};
```

- ▶ Imagine that **cacheValid** is **false**, and then:
 - ▶ One thread calls **Widget::magicValue** and executes through the point where **cacheValid** is set to **true**.
 - ▶ At that moment, a second thread calls **Widget::magicValue** and checks **cacheValid**. Seeing it true, the thread returns **cachedValue**, even though the first thread has not yet made an assignment to it. The returned value is therefore incorrect.

```

class Widget {
public:
    ...

    int magicValue() const
    {
        std::lock_guard<std::mutex> guard(m);    // lock m

        if (cacheValid) return cachedValue;
        else {
            auto val1 = expensiveComputation1();
            auto val2 = expensiveComputation2();
            cachedValue = val1 + val2;
            cacheValid = true;
            return cachedValue;
        }
    }
    // unlock m

    ...

private:
    mutable std::mutex m;
    mutable int cachedValue;                // no longer atomic
    mutable bool cacheValid{ false };        // no longer atomic
};

```

- For a single variable or memory location requiring synchronization, use of a **std::atomic** is adequate.
- but once you get to two or more variables or memory locations that require manipulation as a unit, you should reach for a **mutex**.

EASY PEASY



LEMON SQUEEZY

makeameme.org